

## OS

- Purpose: Hardware Abstraction; Resource Management (Time & Space Multiplexing)

### User Mode VS Kernel Mode

- *User Mode*: cannot issue privileged instructions; can access parts of memory allowed by kernel mode code
- *Kernel Mode*: can issue all instructions; access all memory
- *Privileged*: instructions/memory locations that can interfere with other processes

### I/O Devices

- Device → Controller → Driver
- Accessed via system calls (I/O calls are privileged instructions)
- Can be done using:
  - Busy Waiting: Driver sits in tight loop; poll device (see if done/busy)
  - Interrupts: Device sends interrupt to driver when done
  - Direct Memory Access (DMA): chip controls flow of bits between memory and controller without CPU intervention; chip causes interrupt when done

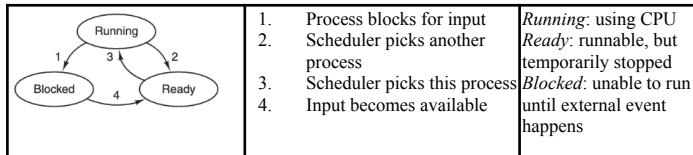
### System Calls

- *System Call*: allows user mode program/process to request a service from OS kernel
  - switches user→kernel mode; significant performance overhead
  - Process, File, Directory Management; kill process, get time
- 1. Put the system call number in a place where OS expects it (register)
- 2. Execute a TRAP instruction to switch from user mode to kernel mode and start execution at a fixed address within the kernel
- 3. The kernel code that starts following the TRAP examines the system-call number and dispatches to the correct system-call handle
- 4. System call handler runs
- 5. Once it has completed its work, control may be returned to the user program
- 6. The system call may block the caller, then the OS will look around to see if some other process can be run next

## Inter-Process Communication

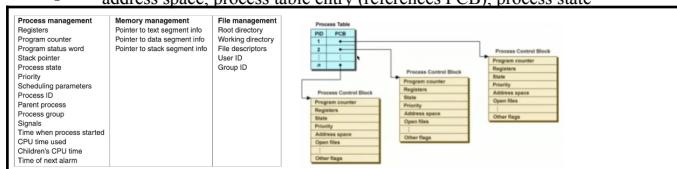
### Process

- *Process*: abstraction of a running program; encapsulates all resources/information needed to run a program



### Process Table

- linked list of PCBs (Process Control Blocks)
  - *PCB*: contains state required to resume a process
  - ICurrent process runs on RAM
  - Not-currently-running process: info stored in process table
  - address space, process table entry (references PCB), process state



### Threads

- *Thread*: sequential execution stream within the process; entities scheduled for execution by CPU
- *PRO*: parallel; shared address space; reduced overhead (easier create, destroy, switch, communicate);  I/O
- *CION*: not help performance when CPU bound (more threads than CPU cores)

Per Process Items (shared among threads)	Per Thread
- Address space	- Program counter
- Global Variables	- Registers
- Open files	- Stack
- Child processes	- State
- Pending alarms	
- Signals and signal handlers	Threads can read, write or delete other thread's stacks
- Accounting information	

### Multiprogramming

- allows multiple programs to share CPU and run in (pseudo)parallel
- but CPU can only run one program at a time – OS gives illusion of parallelism by switching between processes quickly
- processes are sequential but execution is interleaved
- programs cannot rely on timing assumptions

### Race Conditions

- when multiple processes/threads read/write from common storage

stack //global shared variable	Avoiding Race Conditions
<pre>void my_stack_function() {     A if (isEmpty(stack)) return;     B int s = pop(stack);     //do something with s... }</pre>	<ul style="list-style-type: none"> <li>- <i>Mutual Exclusion</i>: ensure that 1 process exists in the critical region at any given time</li> <li>- <i>Critical Region (CR)</i>: part of program where shared resource/data is accessed</li> <li>- Conditions to achieve good solution           <ol style="list-style-type: none"> <li>1. No two processes may be <u>simultaneously</u> inside their CRs</li> <li>2. No <u>assumptions</u> can be made about <u>speed and number of CPUs</u></li> </ol> </li> </ul>
<b>Another potential execution transcript</b> <ol style="list-style-type: none"> <li>1. A tests isEmpty(stack)⇒ false</li> <li>2. B tests isEmpty(stack)⇒ false</li> <li>3. A pops ⇒ stack is now empty</li> <li>4. B pops – Exception?</li> </ol>	

	<ul style="list-style-type: none"> <li>3. No process running <u>outside</u> its CR can <u>block</u> other processes</li> <li>4. No process should have to <u>wait forever</u> to enter CR</li> </ul>
--	--

### Deadlocks

- 4 Conditions for deadlock

1. Mutual Exclusion: each resource can only be assigned to 1 process at a time
2. Hold-and-Wait: process currently holding resources can request new resources
3. No Preemption: resources previously granted cannot be forcibly taken away
4. Circular Wait: circular list of 2 or more processes; each waiting for resource held by next member of chain

### Dealing with Deadlocks

1. Ignore the problem
2. Detection (Resource allocation graph and graph algorithms) and recovery
3. Avoid by careful resource allocation
4. Prevention

### Process Scheduling

- *Scheduling Algorithm*: determines which process to run next
- General Scheduling Goals
- *Fairness*: giving each process a fair share of the CPU
- *Policy Enforcement*: seeing that stated policy is carried out
- *Balance*: keeping all parts of the system busy
- When to schedule: process created/exits/blocks, I/O interrupt, clock interrupt

### Preemptive VS Non-preemptive Scheduling

- *Preemptive*: pick a process and let it run for a quantum (maximum fixed amount of time); clock interrupt at the end of time interval to give control to CPU back to scheduler
- *Non-preemptive*: pick process and let it run for as long as it wants until blocks/voluntarily releases to CPU
- Context Switching: happens every time OS switches process; expensive

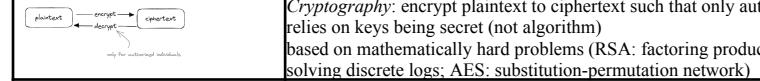
### Scheduling Environments

Batch	Interactive	Real Time
-Don't require user interaction -A set of (potentially periodic) jobs that need to run	-Interactive users expect fast response times/responsiveness from the system	-Need to meet deadlines -Not really focused in this subject
Can be Preemptive or Non-Preemptive	<ul style="list-style-type: none"> <li>Preemptive</li> <li>-by interrupting the process after a fixed time, it ensures that every process will produce a fast response</li> <li>-rather than having to wait, a very long time for one process to finish before moving onto the next one, which may lead to the user having to wait and thus not achieving fast interactivity</li> </ul>	<ul style="list-style-type: none"> <li>Can be Preemptive or Non-Preemptive</li> <li>-possibly schedule the process with fastest deadline first</li> </ul>
	<ul style="list-style-type: none"> <li>-Throughput: maximise jobs per hour</li> <li>-Turnaround time: minimise time between submission and termination</li> <li>-CPU utilisation: keep CPU busy all the time</li> </ul>	<ul style="list-style-type: none"> <li>-Response Time: respond to requests quickly</li> <li>-Proportionality: meet users' expectations</li> </ul>
-[Non-Pre] First-Come First Served (FCFS) -[Non-Pre] Shortest Job First (SJF) -[Pre] Shortest Remaining Time Next (SRT)	[Pre] Round Robin	[Pre] Priority Scheduling

### Memory Management

- *PRO*: support multiprogramming; security (isolation between process and OS memory); enable running processes with large memory requirements
- *Memory Manager*: manages (part of) memory hierarchy; keep track of which parts of memory are in use; allocate when needed – deallocate when done
- *Swapping*: bringing in each process in its entirety, running it for a while

### Cryptography



### CIA

- *Confidentiality*: only sender and intended receiver should understand contents of transmitted message
- *Authentication*: establish identities of one or both endpoints
- *Integrity*: ensure that messages are not altered

### Symmetric Cryptography

- Same key used for encryption and decryption
- *Advanced Encryption Standard (AES)*
- Break data into blocks and encrypt each block
- Modes of operation: ECB, CBC – determines how each block is treated/linked
- *Electronic Codebook (ECB)*





**Congestion Window (CWND)**

- additional window dynamically adjusted
- based on **network performance**
- controlled by the sender
  - knows more information: what is sent and what is received
  - how it works
    1. start with small segment size
    2. if get full window of acknowledgements
      - double congestion window size
    3. grows until
      - timeout
      - threshold (Ssthresh) reached

TCP sender      TCP receiver

Acknowledgement

cwnd = 1      Data

cwnd = 2      1 RTT, 1 packet

cwnd = 3      1 RTT, 2 packets

cwnd = 5      1 RTT, 4 packets

cwnd = 6      1 RTT, 4 packets (pipe is full)

TN 6th 6-44

4. If packet loss  
Ssthresh = CWND/2, start slow start again

**Incremental Congestion Control (Additive Increase)**

5. once Ssthresh is reached
  - slow down growth
  - Additive increase: CWND += 1MSS for each complete window
  - until packet loss (4.)

TN 6th 6-44

**Fast Recovery**

4. if packet loss
  - start from new Ssthresh instead of original starting value
  - avoid slow start phase and go straight to Incremental Congestion Control

TN 6th 6-44

**SACK (Selective Acknowledgements)**

- normal TCP packet loss
  - receiver can only ACK last contiguous packet
  - sender has to retransmit all packets from that point onward
- SACK packet loss
  - TCP header is a little longer
  - receiver can ACK non-contiguous packets
  - receiver can indicate precisely which packets were successful
    - up to 3 ranges of bytes can be specified (TCP header is still limited)
  - sender only retransmits specific missing packets
- avoid unnecessary retransmissions and timeouts

**Retransmit 2 and 5**

Sender      Lost packets      Receiver

ACK: 1      ACK: 1      SACK: 3

ACK: 1      SACK: 3-4      SACK: 6, 3-4

**Windows**

- W increases once per window
  - Acknowledgment back-ack that arrives increases W by 1/W
  - When a loss occurs (with probability p), W is halved
    - with probability p, W <= W/2
  - The average increase in window size is
 
$$\frac{(1-p)}{p} = \frac{W}{1-p}$$
  - To be in equilibrium (balance), the average increase must be 0
 
$$W = \sqrt{2/p}$$

more losses decreases window size

**Rate and Fairness**

- The window is sent  $\leq$  once per round trip time (RTT), T
- Rate is  $\frac{W}{T} = \frac{1}{\sqrt{2/p}}$
1. For a given packet loss rate, longer RTTs get less rate
  - "RTT unfairness"
2. If RTT is small, TCP forces the packet loss rate to be high
  - This formula is very approximate
    - Window only responds to one packet loss per RTT
    - Packet losses are clustered